

## DS8003 – Final Project Report

(By:- Akshdeep Kaler: 501009212, Li Gong: 501077874)

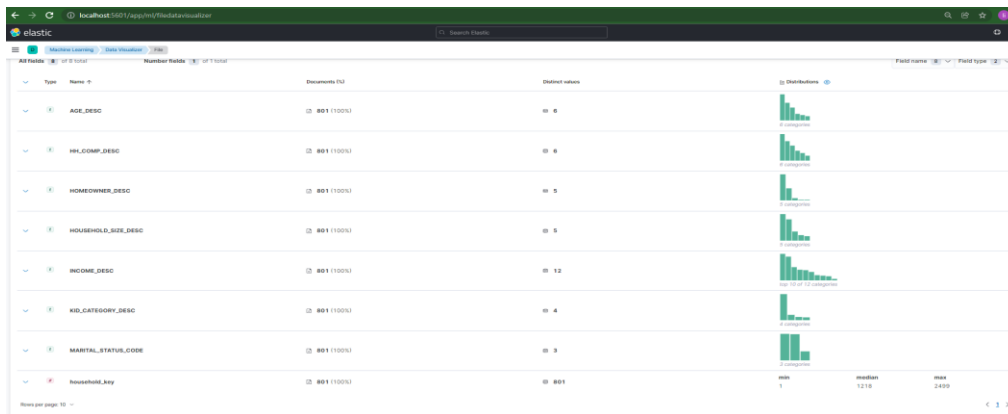
### Problem Definition:

The retail stores sell products to customers and they would like to retain their customers and make them buy more products for increased profitability. For this purpose, various marketing campaigns are also run, in addition to improved service delivery and better pricing. To know if the employed strategy is working, it is important to know that what elements of the strategy steer the business in right directions and which of them are not producing the desired output.

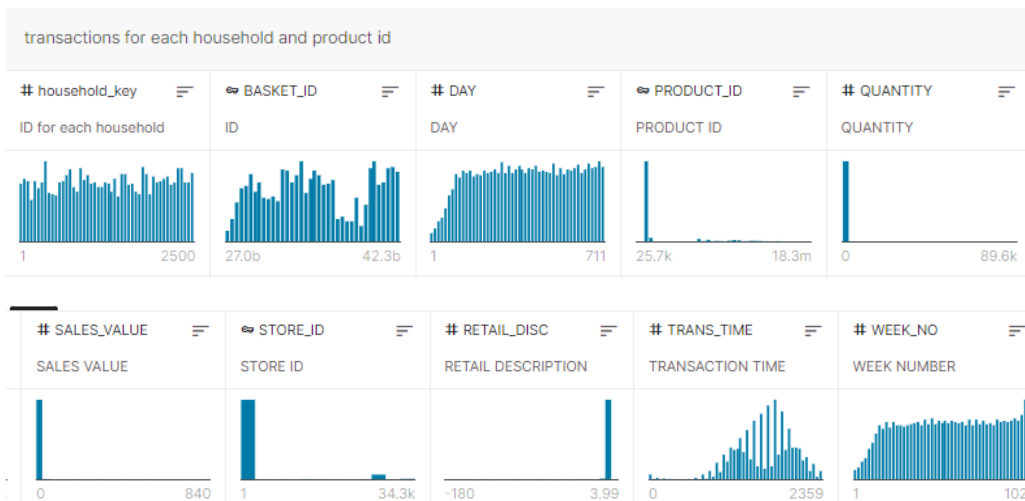
### The Dataset:

We use the open-sourced dataset, titled The Complete Journey, made available by dunhumby. The dataset contains household level, anonymized, transaction data (2500 households) including the demographics and marketing campaigns (30 campaigns). The transactional data include over 90 thousand products categorized in 44 departments. The dataset is available at <https://www.kaggle.com/frtgenn/dunhumby-the-complete-journey>. The following is the summary of each datasets:

1. hh\_demographic: The table contains demographic information for a portion of households.



2. transaction\_data: This contains all products purchased by households.



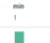

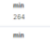
### 3. campaign\_table (Not Used During Project)

This table lists the campaigns received by each household in the study.

4. campaign\_desc: This table gives the length of time for which a campaign runs. Any coupons received as part of a campaign are valid within dates contained in this table.

File stats

All fields: 4 of 4 total      Number fields: 3 of 3 total      Field name:      Field type:





Type	Name	Documents (%)	Distinct values	Distributions
✓	CAMPAIGN	30 (100%)	30	 min: 1, median: 15.5, max: 30
✓	DESCRIPTION	30 (100%)	3	 min: 264, median: 502, max: 719
✓	END_DAY	30 (100%)	28	 min: 224, median: 475, max: 659
✓	START_DAY	30 (100%)	27	

Rows per page: 10

5. product: This table contains information on each product sold such as type of product, national or private label and a brand identifier.

File stats

All fields: 7 of 7 total      Number fields: 3 of 3 total      Field name:      Field type:

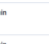


Type	Name	Documents (%)	Distinct values	Distributions
✓	BRAND	999 (100%)	2	 min: 2, median: 69, max: 1820
✓	COMMODITY_DESC	999 (100%)	158	 min: 2, median: 69, max: 1820
✓	CURR_SIZE_OF_PRODUCT	999 (100%)	287	 min: 2, median: 69, max: 1820
✓	DEPARTMENT	999 (100%)	11	 min: 2, median: 69, max: 1820
✓	MANUFACTURER	999 (100%)	176	
✓	PRODUCT_ID	999 (100%)	999	
✓	SUB_COMMODITY_DESC	999 (100%)	301	

Rows per page: 10

6. coupon: This table list all the coupons sent to customers as part of a campaign, as well as the products for which each coupon is redeemable.

File stats

All fields: 3 of 3 total      Number fields: 3 of 3 total      Field name:      Field type:





Type	Name	Documents (%)	Distinct values	Distributions
✓	CAMPAIGN	999 (100%)	23	 min: 1, median: 22, max: 29
✓	COUPON_UPC	999 (100%)	152	 min: 10000085378, median: 53620010028, max: 58768484040
✓	PRODUCT_ID	999 (100%)	469	 min: 27160, median: 85737, max: 161406

Rows per page: 10

7. coupon\_redempt: This table identifies the coupons that each household redeemed.

File stats

All fields: 4 of 4 total      Number fields: 4 of 4 total      Field name:      Field type:

Type	Name	Documents (%)	Distinct values	Distributions
✓	CAMPAIGN	999 (100%)	25	 min: 1, median: 14, max: 30
✓	COUPON_UPC	999 (100%)	378	 min: 10000085320, median: 52370020076, max: 58978500076
✓	DAY	999 (100%)	243	 min: 227, median: 533, max: 691
✓	household_key	999 (100%)	193	 min: 1, median: 472, max: 1098

Rows per page: 10

8. casual\_data (Not Used During Project)

This table signifies whether a given product was featured in the weekly mailer or was part of an in-store display (other than regular product placement).

### **Work Distribution:**

To get the solution to the problem mentioned above, the questions were divided among the group members:

1. Akshdeep Kaler:

- What are the characteristics of customers whose spending at the store is increasing?
- What are the categories of the products that are seeing increased/decreased sales?
- What are the most profitable categories of the products over time?
- Which day has the highest sales?

2. Li Gong:

- Are the marketing campaigns effective?
- Which of the marketing campaigns was the most successful one?
- What are the characteristics of customers who were attracted by each marketing campaign?

Q1. What are the characteristics of customers whose spending at the store is increasing?

Dataset : **transaction\_data.csv** and **hh\_demographic.csv** are used.

The tools used for extracting the information are:

- Pyspark
- Pyspark SQL
- Hadoop Distributed File System (HDFS)

In the data file 'transaction\_data.csv,' each customer or household details of spending is given. For the required question, we are interested in household\_key and sales attributes/columns. The Pyspark API is used to extract the columns and store them on the Hadoop Distributed File System. The following code extracts the total sales according to each 'household\_key' and stores in the HDFS, where it can be accessed for further analysis.

```

from pyspark import SparkConf, SparkContext

def avg(inp):
    sal = 0
    for i in inp[1]:
        sal += float(i)

    a = [inp[0], str(sal)]
    b = ",".join(a)
    return b

def main(sc):
    sc = SparkContext.getOrCreate(SparkConf())
    textFile = sc.textFile("/user/root/FinalProject/datasets/transaction_data.csv")
    header = textFile.first()
    tags = sc.parallelize([header])
    data = textFile.subtract(tags)
    cus_r = data.map(lambda line: line.split(','))
    cus_1 = cus_r.map(lambda field: (field[0],field[5]))
    cus_2 = cus_1.groupByKey()
    cus_3 = cus_2.map(lambda inp: avg(inp))
    cus_4 = cus_3.filter(lambda out: out != None)
    cus_4.coalesce(1).saveAsTextFile("/user/root/FinalProject/q1_output" )

if __name__ == "__main__":
    conf = SparkConf().setAppName("Testing Spark Commands")
    sc = SparkContext(conf = conf)
    main(sc)
    sc.stop()

```

Execution of 'q\_1.py' file in the command shell.

```

[root@sandbox-hdp FinalProject]# ls
datasets q_1.py
[root@sandbox-hdp FinalProject]# spark-submit --master yarn-client --executor-memory 512m --num-executors 3 --executor-cores 1 --driver-memory q_1.py

```

Preview of extracted file in a format where the first tuple is household key and the second is total money spent in-store during two-year data.

```

[root@sandbox-hdp FinalProject]# hadoop fs -ls /user/root/FinalProject/q1_output
Found 2 items
-rw-r--r-- 1 root hdfs 0 2021-11-11 22:55 /user/root/FinalProject/q1_output/_SUCCESS
-rw-r--r-- 1 root hdfs 70966 2021-11-11 22:55 /user/root/FinalProject/q1_output/part-00000
[root@sandbox-hdp FinalProject]# hadoop fs -cat /user/root/FinalProject/q1_output/part-00000 | head -n 10
(u'2068', 2561.3999999999955)
(u'1869', 3630.6199999999979)
(u'1524', 2800.8099999999997)
(u'1942', 3135.0599999999987)
(u'2491', 950.77000000000006)
(u'1946', 3312.4599999999989)
(u'2327', 1267.7600000000007)
(u'818', 1724.5900000000013)
(u'667', 1121.2000000000001)
(u'340', 1057.5599999999997)
[root@sandbox-hdp FinalProject]#
[root@sandbox-hdp FinalProject]#

```

The avoid the first row of the data file 'hh\_demographic.csv' to be included during the Pyspark SQL implementation. Pyspark API in python file is used to remove the first row or header row and store it in HDFS for further processing.

```

from pyspark import SparkConf, SparkContext

def main(sc):
    sc = SparkContext.getOrCreate(SparkConf())
    textFile = sc.textFile("/user/root/FinalProject/datasets/hh_demographic.csv")
    header = textFile.first()
    tags = sc.parallelize([header])
    data = textFile.subtract(tags)
    data.coalesce(1).saveAsTextFile("/user/root/FinalProject/q1_a_output")

if __name__ == "__main__":
    conf = SparkConf().setAppName("Testing Spark Commands")
    sc = SparkContext(conf = conf)
    main(sc)
    sc.stop()

```

Execution of above-written python code file 'q\_a\_1.py' in the command shell.

```

[root@sandbox-hdp FinalProject]# spark-submit --master yarn --deploy-mode client --executor-memory 512m --num-executors 3 --executor-cores 1 --driver-memory 512m q_a_1.py

```

The output of file 'q\_a\_1.py' is stored on the Hadoop file distribution system. To check the output, the following commands are used on the command shell.

```

[root@sandbox-hdp FinalProject]# hadoop fs -cat /user/root/FinalProject/q1_a_output/part-00000 | head -10
45-54,A,100-124K,Homeowner,2 Adults No Kids,2,None/Unknown,2407
65+,B,125-149K,Unknown,2 Adults No Kids,2,None/Unknown,2397
45-54,A,50-74K,Homeowner,2 Adults No Kids,2,None/Unknown,1394
45-54,U,35-49K,Homeowner,1 Adult Kids,4,3+,319
35-44,A,35-49K,Homeowner,2 Adults Kids,3,1,1430
35-44,B,25-34K,Unknown,Single Male,1,None/Unknown,2486
35-44,B,35-49K,Renter,1 Adult Kids,2,1,968
35-44,A,75-99K,Homeowner,2 Adults Kids,3,1,574
25-34,B,35-49K,Homeowner,1 Adult Kids,3,2,1226
25-34,U,Under 15K,Unknown,2 Adults Kids,5+,3+,1174

```

The two outputs from the files 'q\_1.py' and 'q\_a\_1.py' are stored on the Hadoop files Distributed system.

Further Pyspark SQL tool is used for extracting the information of spending of each household/customer. The first two tables are created from the extracted files, and then SQL commands are used to inner join the two tables and extract the information in the end. Following are the procedure followed to get information about the customers.

Run the Pyspark on the command shell.

Following commands are used to create a new table 'hh\_demographic' and 'house\_holdsales'.

The output paths from 'q\_1.py' and 'q\_a\_1.py' are used to generate the tables in Pyspark SQL with the help of the following codes:

```
>>> spark.sql("""
... CREATE TABLE hh_demographich (
... AGE_DESC STRING,
... MARITAL_STATUS_CODE STRING,
... INCOME_DESC STRING,
... HOMEOWNER_DESC STRING,
... HH_COMP_DESC STRING,
... HOUSEHOLD_SIZE_DESC STRING,
... KID_CATEGORY_DESC STRING,
... household_key BIGINT)
... USING CSV OPTIONS (path '/user//user/
... ;
... """)
... )
```

```
>>> spark.sql("""
... CREATE TABLE household_sales (
... Household_key int,
... Tot_sale float)
... USING CSV OPTIONS (path '/user/root/FinalProject/q1_output/part-00000')
... """).show()
```

The generated tables are stored in the 'default' database.

```
>>> spark.sql("show tables").show()
+-----+-----+-----+
|database|tableName|isTemporary|
+-----+-----+-----+
| default|employees_bucket|false|
| default| hh_demographich|false|
| default| household_sales|false|
| default| rounds2|false|
| default| test_map|false|
+-----+-----+-----+

>>>
```

The top 3 rows of each table.

```
>>> spark.sql("select * from household_sales").show(3)
+-----+-----+
|Household_key|Tot_sale|
+-----+-----+
|          2068|  2561.4|
|          1869|  3630.62|
|          1524|  2800.81|
+-----+-----+
only showing top 3 rows
```

```
>>> spark.sql("select * from hh_demographich").show(3)
+-----+-----+-----+-----+-----+-----+-----+-----+
|AGE_DESC|MARITAL_STATUS_CODE|INCOME_DESC|HOMEOWNER_DESC|HH_COMP_DESC|HOUSEHOLD_SIZE_DESC|KID_CATEGORY_DESC|household_key|
+-----+-----+-----+-----+-----+-----+-----+-----+
|  45-54|                A|  100-124K|    Homeowner|2 Adults No Kids|                2|    None/Unknown|        2407|
|   65+|                B|  125-149K|     Unknown|2 Adults No Kids|                2|    None/Unknown|        2397|
|  45-54|                A|   50-74K|    Homeowner|2 Adults No Kids|                2|    None/Unknown|        1394|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 3 rows
```

DataFrames are generated from the two tables and stored under separate variables.

```
>>> df1 = sqlContext.sql("SELECT * FROM hh_demographich")
>>> df2 = sqlContext.sql("SELECT * FROM household_sales")
>>> df1.show(2)
+-----+-----+-----+-----+-----+-----+-----+-----+
|AGE_DESC|MARITAL_STATUS_CODE|INCOME_DESC|HOMEOWNER_DESC|HH_COMP_DESC|HOUSEHOLD_SIZE_DESC|KID_CATEGORY_DESC|household_key|
+-----+-----+-----+-----+-----+-----+-----+-----+
|  45-54|                A|  100-124K|    Homeowner|2 Adults No Kids|                2|    None/Unknown|        2407|
|   65+|                B|  125-149K|     Unknown|2 Adults No Kids|                2|    None/Unknown|        2397|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 2 rows

>>> df2.show(2)
+-----+-----+
|Household_key|Tot_sale|
+-----+-----+
|          2068|  2561.4|
|          1869|  3630.62|
+-----+-----+
```

Pyspark join function combines the two tables on the common column of Household\_key and stores them in the final data frame 'fu\_df.'

```
>>> fu_df = df2.join(df1, df2.Household_key == df1.household_key, how = 'inner')
>>> fu_df.show(2)
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Household_key|Tot_sale|AGE_DESC|MARITAL_STATUS_CODE|INCOME_DESC|HOMEOWNER_DESC|HH_COMP_DESC|HOUSEHOLD_SIZE_DESC|KID_CATEGORY_DESC|household_key|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|          2407|  6665.22|  45-54|                A|  100-124K|    Homeowner|2 Adults No Kids|                2|    None/Unknown|        2407|
|          2397|  1095.1|   65+|                B|  125-149K|     Unknown|2 Adults No Kids|                2|    None/Unknown|        2397|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 2 rows
```

The columns in the final data frame are sorted according to 'Tot\_sale' or total sales done by each customer/ household in descending order. The top 10 customer demographic details are extracted.

```
>>> fu_df.sort(fu_df.Tot_sale.desc()).show(10)
```

Household_key	Tot_sale	AGE_DESC	MARITAL_STATUS_CODE	INCOME_DESC	HOMEOWNER_DESC	HH_COMP_DESC	HOUSEHOLD_SIZE_DESC	KID_CATEGORY_DESC	household_key
1609	27859.68	45-54	A	125-149K	Homeowner	2 Adults Kids	5+	3+	1609
2322	23646.92	45-54	U	175-199K	Homeowner	Single Male	1	None/Unknown	2322
1453	21661.29	45-54	A	125-149K	Homeowner	2 Adults Kids	3	1	1453
1430	20352.99	35-44	A	35-49K	Homeowner	2 Adults Kids	3	1	1430
718	19299.86	45-54	A	25-34K	Homeowner	2 Adults Kids	5+	3+	718
707	19194.42	25-34	A	100-124K	Homeowner	2 Adults Kids	5+	3+	707
1653	19153.75	35-44	B	Under 15K	Homeowner	Single Female	1	None/Unknown	1653
982	18790.34	45-54	U	35-49K	Unknown	2 Adults Kids	4	2	982
400	18494.14	35-44	A	150-174K	Homeowner	2 Adults Kids	3	1	400
1229	18304.31	55-64	A	150-174K	Homeowner	2 Adults No Kids	2	None/Unknown	1229

```
only showing top 10 rows
>>>
```

The final output concluded that among the top 10 customers whose spending is increasing are 45-54 years of age group. Almost everyone is Homeowner and mainly earns above 100K dollars annually.

Q2. What are the categories of the products that are seeing increased/ decreased sales?

The tools used to answer the question are:

- Hive (for querying)
- Apache Hadoop for storage

The data files used for extracting the information are 'transaction\_data.csv' and 'coupon.csv'. First data is stored at Hadoop Distributed File System. Hive is opened by passing command 'hive' on command shell. Then following procedure is followed:

1. Create Database 'finalproject' to store generated tables.

```
Time taken: 1.218 seconds, Fetched: 5 row(s)
hive> create database finalproject;
OK
```

```
hive> show databases;
OK
advanced_hive
assignment3
default
finalproject
foodmart
twitter
Time taken: 0.016 seconds, Fetched: 6 row(s)
```

```
Time taken: 1.17 seconds, Fetched: 6 row(s)
hive> use finalproject;
OK
```



2. Create the tables 'transaction\_data' and 'product' then load the data into table from the HDFS.

```
hive> use finalproject;
OK
Time taken: 0.238 seconds
hive> create table finalproject.transaction_data (
  > household_key string,
  > basket_id string,
  > day string,
  > product_id string,
  > quantity int,
  > sales_value float,
  > store_id string,
  > retail_disc float,
  > trans_time string,
  > week_no string,
  > coupon_disc float,
  > coupon_match_disc float)
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY ',';
OK
Time taken: 0.869 seconds
```

```
hive> create table product (
  > product_id string,
  > manufacturer string,
  > department string,
  > brand string,
  > commodity string,
  > sub_commodity_desc string,
  > size_of_product string)
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY ',';
OK
Time taken: 0.607 seconds
```

```
hive> load data inpath '/user/root/FinalProject/datasets/product.csv'
  > overwrite into table finalproject.product;
Loading data to table finalproject.product
chgrp: changing ownership of 'hdfs://sandbox-hdp.hortonworks.com:8020/apps/hive/warehouse/finalproject.db/product/product.csv': User null does not belong to h
adoop
Table finalproject.product stats: [numFiles=1, numRows=0, totalSize=6429896, rawDataSize=0]
OK
Time taken: 1.192 seconds
hive> show tables;
OK
product
transaction_data
Time taken: 0.24 seconds, Fetched: 2 row(s)
hive> select * from finalproject.product limit 5;
OK
PRODUCT_ID  MANUFACTURER  DEPARTMENT  BRAND  COMMODITY_DESC  SUB_COMMODITY_DESC  CURR_SIZE_OF_PRODUCT
25671  2  GROCERY National  FRZN ICE  ICE - CRUSHED/CUBED  22 LB
26081  2  MISC. TRANS.  National  NO COMMODITY DESCRIPTION  NO SUBCOMMODITY DESCRIPTION
26093  69  PASTRY Private BREAD  BREAD:ITALIAN/FRENCH
26190  69  GROCERY Private FRUIT - SHELF STABLE  APPLE SAUCE  50 OZ
Time taken: 0.135 seconds, Fetched: 5 row(s)
```

### 3. Remove the first row of the tables.

```
hive> ALTER TABLE finalproject.transaction_data
> SET TBLPROPERTIES ("skip.header.line.count"="1");
OK
Time taken: 0.698 seconds
hive> select * from finalproject.transaction_data limit 5;
OK
2375 26984851472 1 1004906 1 1.39 364 -0.6 1631 1 0.0 0.0
2375 26984851472 1 1033142 1 0.82 364 0.0 1631 1 0.0 0.0
2375 26984851472 1 1036325 1 0.99 364 -0.3 1631 1 0.0 0.0
2375 26984851472 1 1082185 1 1.21 364 0.0 1631 1 0.0 0.0
2375 26984851472 1 8160430 1 1.5 364 -0.39 1631 1 0.0 0.0
Time taken: 0.135 seconds, Fetched: 5 row(s)
```

```
hive> ALTER TABLE finalproject.product
> SET TBLPROPERTIES ("skip.header.line.count"="1");
OK
Time taken: 0.707 seconds
hive> select * from finalproject.product limit 5;
OK
25671 2 GROCERY National FRZN ICE ICE - CRUSHED/CUBED 22 LB
26081 2 MISC. TRANS. National NO COMMODITY DESCRIPTION NO SUBCOMMODITY DESCRIPTION
26093 69 PASTRY Private BREAD BREAD:ITALIAN/FRENCH
26190 69 GROCERY Private FRUIT - SHELF STABLE APPLE SAUCE 50 OZ
26355 69 GROCERY Private COOKIES/CONES SPECIALTY COOKIES 14 OZ
Time taken: 0.131 seconds, Fetched: 5 row(s)
```

### 4. Create new table 'dept\_sales' by inner joining the tables 'transaction\_data' and 'product'.

```
hive> CREATE TABLE dept_sales AS
> SELECT p.product_id, p.department, t.sales_value
> FROM product p
> INNER JOIN transaction_data t ON p.product_id = t.product_id;
Query ID = root_20211121183002_2b712ef9-e859-4280-9103-35b9b891e3f6
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1637510490690_0004)
```

	VERTICES	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	.....	SUCCEEDED	1	1	0	0	0	0
Map 2	.....	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 12.15 s
Moving data to directory hdfs://sandbox-hdp.hortonworks.com:8020/apps/hive/warehouse/finalproject.db/dept_sales
Table finalproject.dept_sales stats: [numFiles=1, numRows=2595732, totalSize=52973525, rawDataSize=50377793]
OK
Time taken: 20.122 seconds
hive> describe dept_sales;
OK
product_id      string
department      string
sales_value     float
Time taken: 0.487 seconds, Fetched: 3 row(s)
hive> select * from dept_sales limit 5;
OK
1004906 PRODUCE 1.39
1033142 PRODUCE 0.82
1036325 PRODUCE 0.99
1082185 PRODUCE 1.21
8160430 PRODUCE 1.5
Time taken: 0.134 seconds, Fetched: 5 row(s)
hive>
```

5. Outputs are generated by summing the sales value of each item category in descending and ascending order.

```
hive> SELECT department, round(sum(sales_value),2) sales
> FROM dept_sales
> GROUP BY department
> ORDER BY sales asc
> Limit 10;
Query ID = root_20211121184641_668715fe-fc91-4dc7-997a-403fcab187ee
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637510490690_0004)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED    4         4          0         0         0         0
Reducer 2 ..... SUCCEEDED    1         1          0         0         0         0
Reducer 3 ..... SUCCEEDED    1         1          0         0         0         0
-----
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 6.19 s
-----
OK
      0.0
ELECT & PLUMBING 1.0
GRO BAKERY      2.18
HOUSEWARES     2.99
MEAT-WHSE      7.0
PROD-WHS SALES 7.52
CHARITABLE CONT 7.74
HBC            9.42
TOYS           9.84
PORK           15.7
Time taken: 6.751 seconds, Fetched: 10 row(s)
hive> |
```

Figure 1.

In figure 1, the output generated the sales of bottom 10 categories which had the lowest sales. These categories Elect&Plumbing, Gro Bakery, Housewares, Meat-WHSE, Prod-WHS Sales, HBC, Toys and Pork.

```
hive> SELECT department, round(sum(sales_value),2) sales
> FROM dept_sales
> GROUP BY department
> ORDER BY sales desc
> Limit 10;
Query ID = root_20211121184806_b0dd3ca0-04b9-4e5e-b0e3-dd5ae286c281
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637510490690_0004)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED    4         4          0         0         0         0
Reducer 2 ..... SUCCEEDED    1         1          0         0         0         0
Reducer 3 ..... SUCCEEDED    1         1          0         0         0         0
-----
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 7.32 s
-----
OK
GROCERY 4093814.13
DRUG GM 1055358.02
PRODUCE 557452.11
MEAT    548786.81
KIOSK-GAS 544222.28
MEAT-PCKGD 412436.77
DELI    260866.51
PASTRY  121739.86
MISC SALES TRAN 119960.04
NUTRITION 97669.04
Time taken: 7.906 seconds, Fetched: 10 row(s)
hive>
```

Figure 2.

In figure 2, the output generated shows the top 10 categories which had higher sales. These categories are Grocery, Drug GM, Produce, Meat, Kiosk-Gas, Meat-Pckgd, Deli, Pastry, Misc Sales Tran and Nutrition.

**Q3.** What are the most profitable categories of the products overtime?

From the Figure 2., it is evident the **Groceries** are most profitable category in the stores with sale of around 4,093,814 dollars which is around 50 percent of total sales of the store in 2 years period.

**Q3 a.** Which day in two years period has the highest sales?

The tools used for extracting the information are:

- Hadoop map-reduce (processing)
- HDFS (Hadoop Distributed File System)

The following is procedure followed to get the desired output:

1. Copy the required file data files to the HDFS. In this case, it is 'transaction\_data.csv'.

```
[root@sandbox-hdp datasets]# hadoop fs -put transaction_data.csv /user/root/FinalProject/datasets
[root@sandbox-hdp datasets]# hadoop fs -ls /user/root/FinalProject/datasets
Found 4 items
-rw-r--r--  1 root hdfs      95874 2021-11-11 20:22 /user/root/FinalProject/datasets/campaign_table.csv
-rw-r--r--  1 root hdfs 695858427 2021-11-11 20:22 /user/root/FinalProject/datasets/causal_data.csv
-rw-r--r--  1 root hdfs   44349 2021-11-11 20:22 /user/root/FinalProject/datasets/hh_demographic.csv
-rw-r--r--  1 root hdfs 141742346 2021-11-23 02:06 /user/root/FinalProject/datasets/transaction_data.csv
[root@sandbox-hdp datasets]#
```

2. Created the Mapping python file, in which first row is skipped and output the two columns value in key, value pair. In first part of pair 'word\_1' refers to the 'day' column and in second part 'word\_2' refers to the 'sales\_value' column.

```
#!/usr/bin/env python
# Mapping
import sys

for line in sys.stdin:
    # line = line.strip()
    words = line.split(',')
    head_ls = ['household_key', 'BASKET_ID', 'DAY', 'PRODUCT_ID', 'QUANTITY', 'SALES_VALUE', 'STORE_ID', 'RETAIL_DISC',
               'TRANS_TIME', 'WEEK_NO', 'COUPON_DISC', 'COUPON_MATCH_DISC']
    word_1 = ""
    word_2 = ""
    for word in words:
        if (word not in head_ls):
            if(word == words[2]):
                word_1 = word
            elif(word == words[5]):
                word_2 = word
    print("%s,%s" %(word_1, word_2))
```

3. The reducer python file 'wc\_reducer.py' is created to read through the output of the mapping python file 'wc\_mapper.py' and performs the operations and extract the day which has the highest sales value.

```
#!/usr/bin/env python
# Reducer
import sys

current_word = " "
tot_sal = {}
for line in sys.stdin:
    line = line.strip()
    word1, count = line.split(",")
    try :
        count = float(count)
    except ValueError:
        count = 0

    if word1 != current_word:
        current_word = word1
        tot_sal[current_word] = count
    else:
        tot_sal[current_word] += count

print("The day %s has the highest sale of %s dollars." % (max(tot_sal, key=tot_sal.get), max(tot_sal.values())))
```

4. The mapper and reducer files are copied into the local storage of sandbox.

```
[root@sandbox-hdp lab]# ls
wc_mapper.py  wc_reducer.py
[root@sandbox-hdp lab]#
```

5. Python streaming is executed and it uses wc\_mapper.py and wc\_reducer.py files to get the output.

```
[root@sandbox-hdp lab]# hadoop jar /usr/hdp/2.6.5.0-292/hadoop-mapreduce/hadoop-streaming-2.7.3.2.6.5.0-292.jar -files /root/lab/wc_mapper.py,/root/lab/wc_reducer.py -mapper wc_mapper.py -reducer wc_reducer.py -input /user/root/FinalProject/datasets/transaction_data.csv -output /user/root/FinalProject/q3_output
```

```
21/11/23 02:13:07 INFO mapreduce.Job: Job job_1637624652147_0002 completed successfully
```

6. The output files are generated.

```
21/11/23 02:13:07 INFO streaming.StreamJob: Output directory: /user/root/FinalProject/q3_output
[root@sandbox-hdp lab]# hadoop fs -ls /user/root/FinalProject/q3_output
Found 2 items
-rw-r--r--  1 root root      0 2021-11-23 02:13 /user/root/FinalProject/q3_output/_SUCCESS
-rw-r--r--  1 root root    54 2021-11-23 02:13 /user/root/FinalProject/q3_output/part-00000
[root@sandbox-hdp lab]# hadoop fs -cat /user/root/FinalProject/q3_output/part-00000
The day 641 has the highest sale of 24760.1 dollars.
[root@sandbox-hdp lab]#
```

The output shows that 641th day of 2 year period had the highest sales of 24740.1 dollars. The 641th day falls in the month of October. From the value we infer the reason behind the sales rise.

**Q4:** Are the marketing campaigns effective?

Datasets : **transaction\_data.csv** , **coupon\_redempt.csv**

Tools:

- Hive: create tables, query tables
- Hadoop Distributed File System (HDFS) : store distributed data

Aiming to prove the effects of marketing campaigns, whether the coupons distributed during campaigns can increase the sales by comparing the sales values and quantities between purchase with redeemed coupons and without coupons.

1. Create tables: **transaction\_data**, **coupon\_redempt**

```
hive> load data inpath '/user/root/finalproject/transaction_data.csv' overwrite
into table final.transaction;
Loading data to table final.transaction
chgrp: changing ownership of 'hdfs://sandbox-hdp.hortonworks.com:8020/apps/hive/
warehouse/final.db/transaction/transaction_data.csv': User null does not belong
to hadoop
Table final.transaction stats: [numFiles=1, numRows=0, totalSize=141742346, rawD
ataSize=0]
```

```
hive> select * from transaction limit 10;
OK
household_key  BASKET_ID  DAY  PRODUCT_ID  NULL  NULL  STORE_ID
NULL  TRANS_TIME  WEEK_NO  NULL  NULL
2375  26984851472  1  1004906  1  1.39  364  -0.6  1631  1
0.0  0.0
2375  26984851472  1  1033142  1  0.82  364  0.0  1631  1
0.0  0.0
2375  26984851472  1  1036325  1  0.99  364  -0.3  1631  1
0.0  0.0
2375  26984851472  1  1082185  1  1.21  364  0.0  1631  1
0.0  0.0
2375  26984851472  1  8160430  1  1.5  364  -0.39  1631  1
0.0  0.0
2375  26984851516  1  826249  2  1.98  364  -0.6  1642  1
0.0  0.0
2375  26984851516  1  1043142  1  1.57  364  -0.68  1642  1
0.0  0.0
2375  26984851516  1  1085983  1  2.99  364  -0.4  1642  1
0.0  0.0
2375  26984851516  1  1102651  1  1.89  364  0.0  1642  1
0.0  0.0
Time taken: 0.576 seconds, Fetched: 10 row(s)
```



```
hive> CREATE TABLE final.coupon_redempt(
  > household_key string,
  > day string,
  > coupon_upc string,
  > campaign string)
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY ',';
```

```
OK
Time taken: 0.639 seconds
hive>
```

```
hive>
  > load data inpath '/user/root/finalproject/coupon_redempt.csv' overwrite in
to table final.coupon_redempt;
Loading data to table final.coupon_redempt
chgrp: changing ownership of 'hdfs://sandbox-hdp.hortonworks.com:8020/apps/hive/
warehouse/final.db/coupon_redempt/coupon_redempt.csv': User null does not belong
to hadoop
Table final.coupon_redempt stats: [numFiles=1, numRows=0, totalSize=54108, rawDa
taSize=0]
OK
Time taken: 1.257 seconds
hive>
```

```
hive> select * from coupon_redempt limit 10;
OK
```

household_key	DAY	COUPON_UPC	CAMPAIGN
1	421	10000085364	8
1	421	51700010076	8
1	427	54200000033	8
1	597	10000085476	18
1	597	54200029176	18
8	422	53600000078	8
13	396	53700048182	5
13	424	10000085364	8
13	434	53600000078	8

```
Time taken: 0.155 seconds, Fetched: 10 row(s)
```

## 2. Calculate the total values for the purchase with redeemed coupons

After checking the values in transaction table and coupon\_redempt tables, we found that the “household\_key” is a n-to-n matching key. As “household\_key” in both tables are not unique. We can understand as one household had several purchases ( in transactions table) while the same household had multiple records of using coupons ( in coupon\_redempt table). Therefore, in order not to expand the total sales amount after the join, we add another value - "day" - to the join key to restrict the matching condition.

```
hive> SELECT sum(sales_value) as total_value
> FROM transaction t
> INNER JOIN
> coupon_redempt cr
> ON t.household_key = cr.household_key and t.day = cr.day;
Query ID = root_20211125171316_b7e5f9e1-4e97-41af-8935-9e5625881257
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1637766516916_0019)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED    9         9         0         0         0         0
Map 3 .....  SUCCEEDED    1         1         0         0         0         0
Reducer 2 ..... SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 03/03  [=====>>] 100%  ELAPSED TIME: 16.42 s
-----
OK
272944.45904690586
Time taken: 21.99 seconds, Fetched: 1 row(s)
hive>
```

### 3. Calculate the total quantities for the purchase with redeemed coupons

```
hive> SELECT sum(sales_quantity) as total_quantity
> FROM transaction t
> INNER JOIN
> coupon_redempt cr
> ON t.household_key = cr.household_key and t.day = cr.day;
FAILED: SemanticException [Error 10004]: Line 1:11 Invalid table alias or column
reference 'sales_quantity': (possible column names are: t.household_key, t.bask
et_id, t.day, t.product_id, t.quantity, t.sales_value, t.store_id, t.retail_disc
, t.trans_time, t.week_no, t.coupon_disc, t.coupon_match_disc, cr.household_key,
cr.day, cr.coupon_upc, cr.campaign)
hive> SELECT sum(quantity) as total_quantity
> FROM transaction t
> INNER JOIN
> coupon_redempt cr
> ON t.household_key = cr.household_key and t.day = cr.day;
Query ID = root_20211125171944_7bd238c9-e768-44f7-88a9-36084bb0eedc
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637766516916_0019)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED    9         9         0         0         0         0
Map 3 .....  SUCCEEDED    1         1         0         0         0         0
Reducer 2 ..... SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 03/03  [=====>>] 100%  ELAPSED TIME: 25.30 s
-----
OK
3063932
Time taken: 26.048 seconds, Fetched: 1 row(s)
hive>
```



#### 4. Calculate the total values for the purchase without redeemed coupons

This time we use left join to keep all the purchase records from transaction table. The ones which are not matched by the records from coupon\_redempt table are the purchases without redeeming coupons.

```
hive> SELECT sum(sales_value)
> FROM
> (SELECT t.household_key, t.day, sales_value, quantity, cr.campaign
> FROM transaction t
> LEFT JOIN
> coupon_redempt cr
> ON t.household_key = cr.household_key and t.day = cr.day) tmp
> WHERE tmp.campaign IS NULL;
Query ID = root_20211125172647_a17c6ae0-04e0-4ab5-b628-808089705bcc
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637766516916_0019)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED    9         9         0         0         0         0
Map 3 .....  SUCCEEDED    1         1         0         0         0         0
Reducer 2 ..... SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 31.28 s
-----
OK
7936856.7227408085
Time taken: 31.98 seconds, Fetched: 1 row(s)
```

#### 5. Calculate the total quantities for the purchase without redeemed coupons

```
hive> SELECT sum(quantity)
> FROM
> (SELECT t.household_key, t.day, sales_value, quantity, cr.campaign
> FROM transaction t
> LEFT JOIN
> coupon_redempt cr
> ON t.household_key = cr.household_key and t.day = cr.day) tmp
> WHERE tmp.campaign IS NULL;
Query ID = root_20211125172842_0d56elff-c912-4e8f-a692-b652bd6018f8
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637766516916_0019)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED    9         9         0         0         0         0
Map 3 .....  SUCCEEDED    1         1         0         0         0         0
Reducer 2 ..... SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 16.28 s
-----
OK
259284315
Time taken: 16.984 seconds, Fetched: 1 row(s)
```

6. Use another method to validate the total amounts

Calculate the total sales value by selecting total "sales\_values" from transaction table.

```
hive> select sum(sales_value) from transaction;
Query ID = root_20211125173549_5ca137c5-3556-400c-aell-42aa9da39b83
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637766516916_0019)

-----
      VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED      9          9          0          0          0          0
Reducer 2 .....  SUCCEEDED      1          1          0          0          0          0
-----
VERTICES: 02/02  [=====>>] 100%  ELAPSED TIME: 19.28 s
-----
OK
8057463.0522980355
Time taken: 19.253 seconds, Fetched: 1 row(s)
```

Calculate the total sales value by selecting total "quantities" from transaction table.

```
hive> select sum(quantity) from transaction;
Query ID = root_20211125154552_d56301e8-f31d-40b9-b97a-9640b0f1128e
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637766516916_0014)

-----
      VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED      9          9          0          0          0          0
Reducer 2 .....  SUCCEEDED      1          1          0          0          0          0
-----
VERTICES: 02/02  [=====>>] 100%  ELAPSED TIME: 224.70 s
-----
OK
260685622
Time taken: 226.204 seconds, Fetched: 1 row(s)
```

7. Calculate the promotion rate by sales and quantities:

Sales with redeemed coupons / (sales with redeemed coupons + sales without redeemed coupons):

$$272944 / (272944 + 7936856) * 100 = 3.3\%$$

Sales with redeemed coupons / total sales:

$$272944 / 8057463 * 100 = 3.3\%$$

Quantities with redeemed coupons / (quantities with redeemed coupons + quantities without redeemed coupons):

$$3063932 / (3063932 + 259284315) * 100 = 1.2\%$$

Quantities with redeemed coupons / total quantities:

$3063932/260685622 * 100 = 1.2\%$

The marketing campaign has increased sales by 3.3%, 1,2% of quantities. Obviously, the rate of increase in volume is less than half of the increase in sales. We can conclude that marketing campaigns are more effective on higher price products.

**Q5.** Which of the marketing campaigns was the most successful one?

Datasets : **transaction\_data.csv** , **coupon\_redempt.csv**, **campaign\_desc.csv**

Tools:

- Hive: create tables, query tables
- Hadoop Distributed File System (HDFS) : store distributed data

To evaluate the success of the campaign, we can take clues from the coupons redeemed during the campaign by adding up the total sales and quantities relating the redeemed coupons for each campaign. Then check which one has the largest volume.

1. Create tables: **campaign\_desc**

```
hive> CREATE TABLE campaign_desc(  
  > description string,  
  > Campaign string,  
  > start_day string,  
  > end_day string)  
  > ROW FORMAT DELIMITED  
  > FIELDS TERMINATED BY ',';  
OK  
Time taken: 1.302 seconds
```

```
hive> load data inpath '/user/root/finalproject/campaign_desc.csv' overwrite into  
table campaign_desc;  
Loading data to table default.campaign_desc  
chgrp: changing ownership of 'hdfs://sandbox-hdp.hortonworks.com:8020/apps/hive/  
warehouse/campaign_desc/campaign_desc.csv': User null does not belong to hadoop  
Table default.campaign_desc stats: [numFiles=1, numRows=0, totalSize=540, rawDataSize=0]  
OK  
Time taken: 1.195 seconds
```

```
hive> select * from final.campaign_desc limit 10;  
OK  
DESCRIPTION      CAMPAIGN      START_DAY      END_DAY  
TypeB    24      659      719  
TypeC    15      547      708  
TypeB    25      659      691  
TypeC    20      615      685  
TypeB    23      646      684  
TypeB    21      624      656  
TypeB    22      624      656  
TypeA    18      587      642  
TypeB    19      603      635  
Time taken: 0.248 seconds, Fetched: 10 row(s)  
hive>
```

- Find total sales value/quantities of each campaign by descending order

### Calculate total sales value per campaign by descending order.

After inner joining transaction table and coupon\_redempt table (as described in Q4), the query results are grouped by campaign number to get the total sales per campaign. This is then joined to the campaign\_desc table to link with the campaign information while keeping the whole campaign info from the previous query (right join).

```
hive> SELECT tmp.campaign, total_value, description, start_day, end_day, (end_day-start_day) as last_days
> FROM
> campaign_desc c
> RIGHT JOIN
> (SELECT campaign, sum(sales_value) as total_value
> FROM transaction t
> INNER JOIN
> coupon_redempt cr
> ON t.household_key = cr.household_key and t.day = cr.day
> GROUP BY campaign) tmp
> ON c.campaign = tmp.campaign
> ORDER BY total_value DESC;
```

Query ID = root\_20211125164106\_5a301a75-f590-4365-bdb1-300ellecl3f5

Total jobs = 1

Launching Job 1 out of 1

Status: Running (Executing on YARN cluster with App id application\_1637766516916\_0017)

	VERTICES	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....		SUCCEEDED	9	9	0	0	0	0
Map 4 .....		SUCCEEDED	1	1	0	0	0	0
Map 5 .....		SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....		SUCCEEDED	1	1	0	0	0	0
Reducer 3 .....		SUCCEEDED	1	1	0	0	0	0

VERTICES: 05/05 [=====>>] 100% ELAPSED TIME: 91.83 s

```
OK
18      78480.16967327893      TypeA      587      642      55.0
13      74923.95979427546      TypeA      504      551      47.0
8       39087.329840546474      TypeA      412      460      48.0
26      10501.529964849353      TypeA      224      264      40.0
17      7856.639978066087      TypeB      575      607      32.0
23      7317.879988960922      TypeB      646      684      38.0
22      5956.809986650944      TypeB      624      656      32.0
25      5898.519975185394      TypeB      659      691      32.0
16      5679.249978095293      TypeB      561      593      32.0
30      5550.999983474612      TypeA      323      369      46.0
14      5302.659990489483      TypeC      531      596      65.0
20      4768.52996841073      TypeC      615      685      70.0
9       4418.369988203049      TypeB      435      467      32.0
12      4148.169979020953      TypeB      477      509      32.0
19      3326.8099823594093      TypeB      603      635      32.0
29      1915.9700000882149      TypeB      281      334      53.0
5       1416.6799924075603      TypeB      377      411      34.0
10      1316.7799952700734      TypeB      463      495      32.0
11      1021.3799973353744      TypeB      477      523      46.0
24      971.5299966335297      TypeB      659      719      60.0
7       694.9099982976913      TypeB      398      432      34.0
4       659.3799972236156      TypeB      372      404      32.0
21      650.2699986100197      TypeB      624      656      32.0
2       489.5300007760525      TypeB      351      383      32.0
27      221.6599993109703      TypeC      237      300      63.0
15      133.09000077843666      TypeC      547      708      161.0
28      101.25999891757965      TypeB      259      320      61.0
3       63.62999975681305      TypeC      356      412      56.0
1       60.19999969005585      TypeB      346      383      37.0
6       10.559999942779541      TypeC      393      425      32.0
CAMPAIGN      NULL      DESCRIPTION      START_DAY      END_DAY NULL
Time taken: 93.049 seconds, Fetched: 31 row(s)
hive>
```

According to the total sales value, No.18 campaign has the best performance. Top 4 campaigns are mostly type A which last around 45-50 days. Type B campaigns which have length of around 30 days performed in the medium level. Campaigns longer than 55 days (Type C) are not recommended.

### Calculate total quantities per campaign by descending order.

```

hive> SELECT tmp.campaign, total_quantity, description, start_day, end_day, (end_day-start_day) as last_days
> FROM
> campaign_desc c
> RIGHT JOIN
> (SELECT campaign, sum(quantity) as total_quantity
> FROM transaction t
> INNER JOIN
> coupon_redempt cr
> ON t.household_key = cr.household_key and t.day = cr.day
> GROUP BY campaign) tmp
> ON c.campaign = tmp.campaign
> ORDER BY total_quantity DESC;
Query ID = root_20211125164514_38eb55cf-e08d-4330-a4cd-1dc73c18cbbd
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1637766516916_0017)

-----
VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... SUCCEEDED    9         9         0         0         0         0
Map 4 ..... SUCCEEDED    1         1         0         0         0         0
Map 5 ..... SUCCEEDED    1         1         0         0         0         0
Reducer 2 ..... SUCCEEDED    1         1         0         0         0         0
Reducer 3 ..... SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 05/05  [=====>>] 100%  ELAPSED TIME: 50.34 s
-----
OK
13      868953  TypeA    504      551      47.0
18      844892  TypeA    587      642      55.0
8       397183  TypeA    412      460      48.0
16      140980  TypeB    561      593      32.0
19      133062  TypeB    603      635      32.0
22      129783  TypeB    624      656      32.0
26      91644   TypeA    224      264      40.0
17      79167   TypeB    575      607      32.0
23      63117   TypeB    646      684      38.0
25      51650   TypeB    659      691      32.0
4       35130   TypeB    372      404      32.0
7       34499   TypeB    398      432      34.0
12      30082   TypeB    477      509      32.0
29      29950   TypeB    281      334      53.0
10      29360   TypeB    463      495      32.0
30      28590   TypeA    323      369      46.0
11      19026   TypeB    477      523      46.0
27      15441   TypeC    237      300      63.0
20      13974   TypeC    615      685      70.0
24      13322   TypeB    659      719      60.0
21      8875    TypeB    624      656      32.0
14      2369    TypeC    531      596      65.0
9       1948    TypeB    435      467      32.0
5       473     TypeB    377      411      34.0
2       296     TypeB    351      383      32.0
15      65      TypeC    547      708      161.0
28      40      TypeB    259      320      61.0
1       30      TypeB    346      383      37.0
3       27      TypeC    356      412      56.0
6       4       TypeC    393      425      32.0
CAMPAIGN      NULL  DESCRIPTION  START_DAY      END_DAY NULL
Time taken: 51.573 seconds, Fetched: 31 row(s)

```

As for the total quantities, No. 13 has the largest volume among all the campaign. By comparing the top campaigns between sales and quantities. Some campaigns, e.g. no. 18, 26, 17, 23, 25, are more effective on valuable products (with higher price). As these are also top campaigns, we could conclude products with higher price are sensitive to campaigns.

In conclusion, No. 18 is the most successful campaign due to its largest sales contribution. In addition, it may consist of premium products that are sensitive to the marketing campaign.

**Q6.** What are the characteristics of customers who were attracted by each marketing campaign?

Datasets : **hh\_demographic.csv**, **coupon\_redempt.csv**

Tools:

- Pyspark : create tables, query tables, output csv file
- Hadoop Distributed File System (HDFS) : store distributed data, output query result
- Elasticsearch and Kibana: visualize query result

In order to get the information of customers for each marketing campaign, we can look into the coupon\_redempt table, then link these customers with their demographic info from the hh\_demographic table.

1. Create RDDs in Pyspark:

Create hh\_demographic RDD:

```
>>> hh_demographic = spark.read.options(header=True).csv("/user/root/finalproject/hh_demographic.csv");
```

```
>>> hh_demographic.select("*").show(10);
```

AGE_DESC	MARITAL_STATUS_CODE	INCOME_DESC	HOMEOWNER_DESC	HH_COMP_DESC	HOUSEHOLD_SIZE_DESC	KID_CATEGORY_DESC	household_key
65+	A	35-49K	Homeowner	2 Adults No Kids	2		1
45-54	A	50-74K	Homeowner	2 Adults No Kids	2		7
25-34	U	25-34K	Unknown	2 Adults Kids	3	1	8
25-34	U	75-99K	Homeowner	2 Adults Kids	4	2	13
45-54	B	50-74K	Homeowner	Single Female	1		16
65+	B	Under 15K	Homeowner	2 Adults No Kids	2		17
45-54	A	100-124K	Homeowner	2 Adults No Kids	2		18
35-44	B	15-24K	Unknown	Single Female	1		19
25-34	A	75-99K	Renter	2 Adults No Kids	2		20
45-54	A	75-99K	Homeowner	2 Adults No Kids	2		22

only showing top 10 rows



```
>>> coupon_redempt = spark.read.options(header=True).csv("/user/root/finalproject/coupon_redempt.csv");
>>> coupon_redempt.select("*").show(10);
+-----+-----+-----+-----+
|household_key|DAY| COUPON_UPC|CAMPAIGN|
+-----+-----+-----+-----+
|1|421|10000085364|8|
|1|421|51700010076|8|
|1|427|54200000033|8|
|1|597|10000085476|18|
|1|597|54200029176|18|
|8|422|53600000078|8|
|13|396|53700048182|5|
|13|424|10000085364|8|
|13|434|53600000078|8|
|13|447|52370020076|8|
+-----+-----+-----+-----+
only showing top 10 rows
```

```
>>> hh_demographic.createTempView("hh_demographic");
>>> coupon_redempt.createTempView("coupon_redempt");
```

- To get the information of customers for each campaign, we group the data by campaign number and household ID in coupon\_redempt table. Then join the customer demographic info with keeping all the campaign number from previous query result.

```
>>> result_q6 = spark.sql('''
... SELECT tmp.campaign, d.*
... FROM hh_demographic d
... RIGHT JOIN
... (SELECT campaign,household_key
... FROM coupon_redempt
... GROUP BY campaign,household_key) tmp
... ON d.household_key = tmp.household_key
... ORDER BY tmp.campaign''');
>>> result_q6.select("*").show(20);
```

campaign	AGE_DESC	MARITAL_STATUS_CODE	INCOME_DESC	HOMEOWNER_DESC	HH_COMP_DESC	HOUSEHOLD_SIZE_DESC	KID_CATEGORY_DESC	household_key	
10	1	null	null	null	null	null	null	null	
10	10	45-54	A	25-34K	Homeowner	2 Adults Kids	5+	3+	718
10	10	35-44	A	75-99K	Homeowner	2 Adults Kids	3	1	574
10	10	35-44	U	50-74K	Unknown	Unknown	1	None/Unknown	1541
10	10	45-54	U	35-49K	Unknown	2 Adults Kids	4	2	982
10	10	25-34	A	75-99K	Homeowner	2 Adults Kids	4	2	605
10	10	25-34	U	75-99K	Homeowner	2 Adults Kids	4	2	13
10	10	null	null	null	null	null	null	null	null
10	10	25-34	A	50-74K	Homeowner	2 Adults Kids	4	2	2124
10	10	35-44	U	15-24K	Homeowner	2 Adults Kids	3	1	2280
10	10	35-44	A	150-174K	Homeowner	Single Male	2	None/Unknown	1710
11	11	null	null	null	null	null	null	null	null
11	11	35-44	A	50-74K	Homeowner	2 Adults No Kids	2	None/Unknown	2489
11	11	45-54	U	Under 15K	Homeowner	Single Male	1	None/Unknown	1326
11	11	35-44	B	25-34K	Unknown	Single Female	1	None/Unknown	1707
11	11	45-54	A	125-149K	Homeowner	1 Adult Kids	5+	3+	1451
11	11	25-34	U	75-99K	Homeowner	2 Adults Kids	4	2	13
12	12	55-64	A	150-174K	Homeowner	2 Adults No Kids	2	None/Unknown	1229
12	12	25-34	A	25-34K	Homeowner	2 Adults	3	1	1899
12	12	null	null	null	null	null	null	null	null

only showing top 20 rows

### 3. Output query result in Pyspark

Write the result into csv file.

```
>>> result_q6.write.csv("/user/root/finalproject/result_q6.csv");  
>>> quit();
```

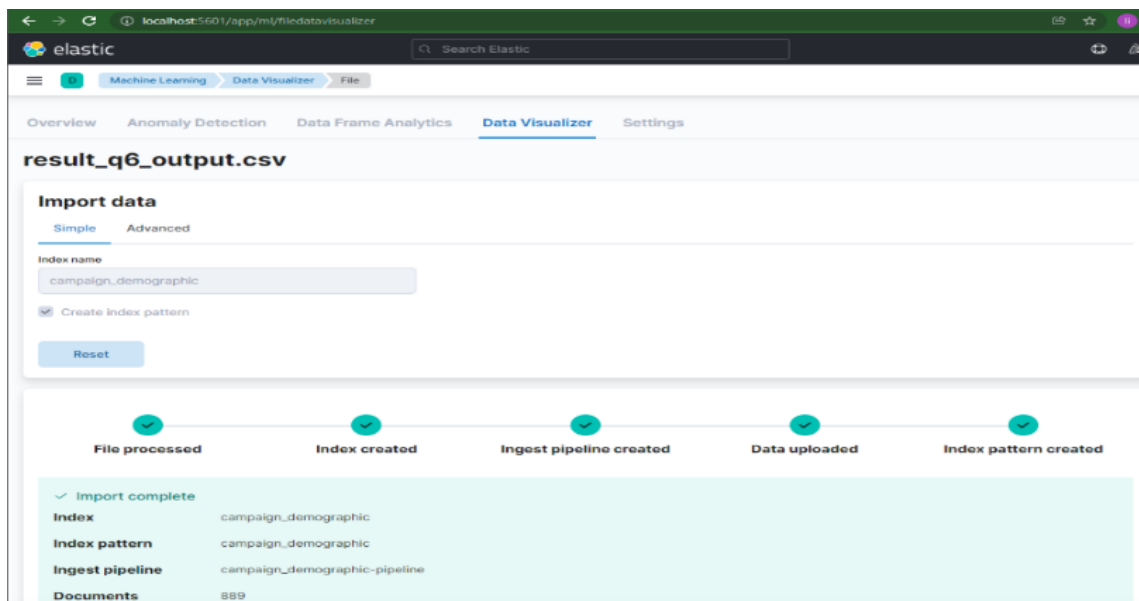
### 4. Output result from HDFS

In order to use csv file in other tool, we merge all the output csv files into one csv file, and output it from HDFS directory to local directory.

```
[root@sandbox-hdp ~]# cd /root/finalproject  
[root@sandbox-hdp finalproject]# hadoop fs -ls /user/root/finalproject  
Found 3 items  
-rw-r--r--    1 root root      54108 2021-11-28 21:27 /user/root/finalproject/coupon_redempt.csv  
-rw-r--r--    1 root root      44349 2021-11-28 21:28 /user/root/finalproject/hh_demographic.csv  
drwxr-xr-x    - root root          0 2021-11-28 22:09 /user/root/finalproject/result_q6.csv  
[root@sandbox-hdp finalproject]# hadoop fs -getmerge /user/root/finalproject/result_q6.csv result_q6_output.csv  
[root@sandbox-hdp finalproject]# ls  
campaign_desc.csv    coupon.csv            product.csv  
campaign_table.csv   coupon_redempt.csv   result_q6_output.csv  
causal_data.csv      hh_demographic.csv   transaction_data.csv  
[root@sandbox-hdp finalproject]#
```

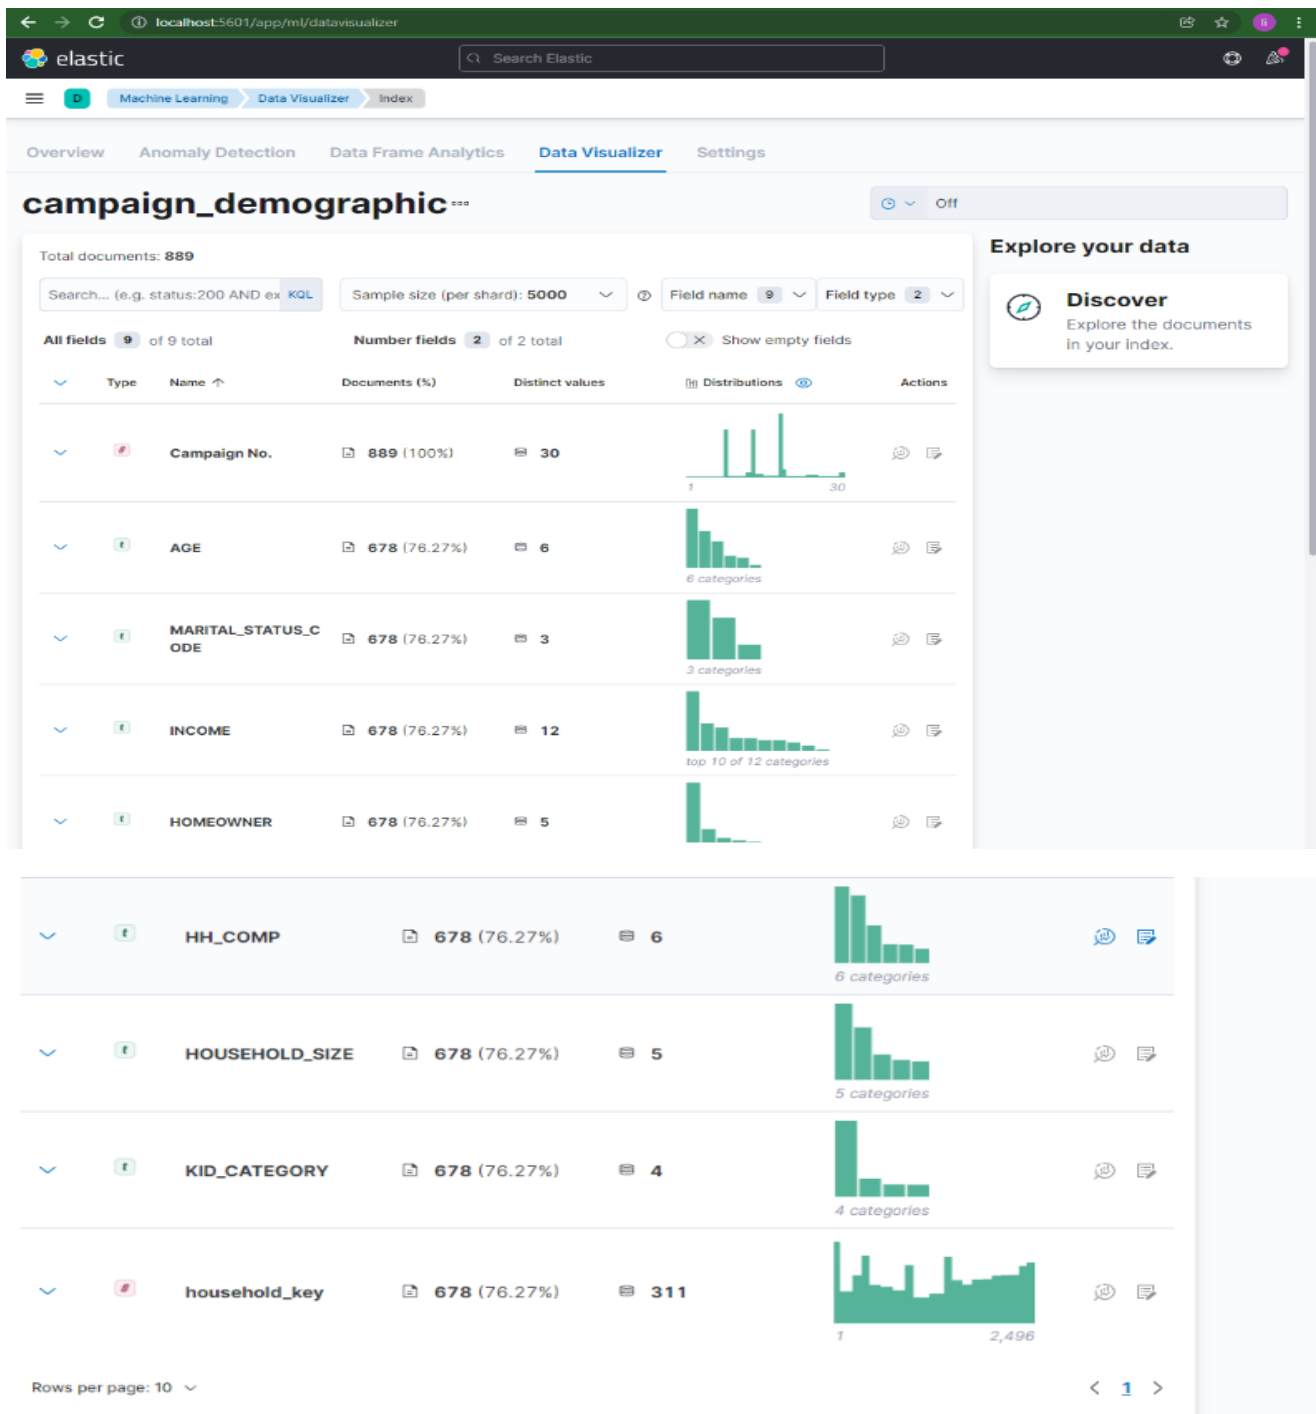
### 5. Data visualization in Kibana

Import query result into Kibana.



Set up Elasticsearch index pattern named campaign\_demographic.

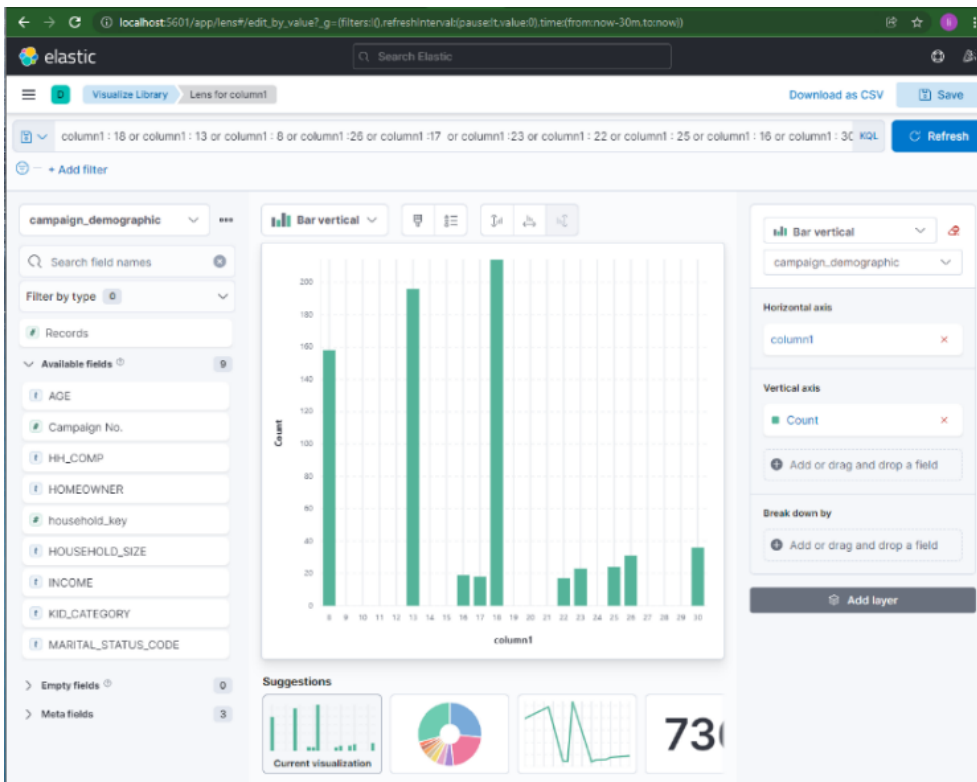




Filter data by doing following search:

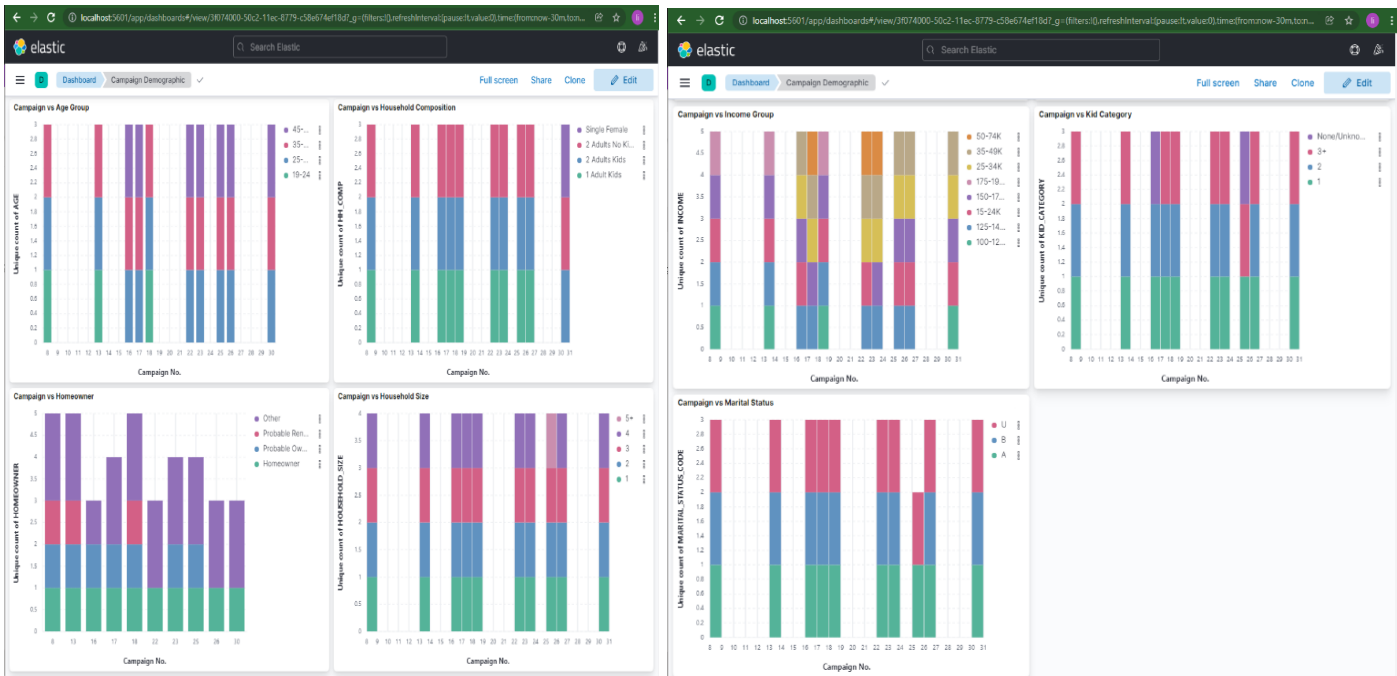
column1: 18 or column1: 13 or column1: 8 or column1 :26 or column1 :17 or column1 :23 or column1:22 or column1: 25 or column1: 16 or column1: 30

These are top 10 campaigns as we found in Q5.



Create a dashboard by illustrating 7 customer characters of top 10 campaigns.

Here is the shared link: <http://localhost:5601/goto/52b4ff7b5e9ca80e8999de952e404d27>



As we can see from the bar chart, most of the campaign fans are between the ages of 45-54, have families consisting of two adults with no children, rent rather than own, and have incomes between \$350,000 and \$490,000, if with more than three children.

### Insights Description:

- Mostly the customers who spend more on the store are of 45-54 year age group have kids and are married. It could be the reason behind the higher sales of the grocery items in the stores.
- As the grocery covers around 50% of total sales and store has \$24,740.1 dollar of highest sale in a day, the store need to originate up with the more ideas to increase the sales of other product.
- Marketing campaigns are more effective on higher price products due to the larger increase of sales than quantities.
- The campaigns with optimal performance mostly last about 45-50 days. We don't recommend more than 55 days.
- Customers between age above 45, with three more children, and income between 35k and 49k are more likely to enrolled in the marketing campaigns.

### Future work:

- We will consider causal dataset to exclude the effects of other occurring events, such as in-store display, weekly mailer, for each campaign.

### References:

- Datasets: <https://www.kaggle.com/frtgnn/dunhumby-the-complete-journey>
- Complete Kibana Tutorial to Visualize and Query Data  
<https://phoenixnap.com/kb/kibana-tutorial>